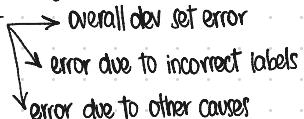


CLEANING UP INCORRECTLY MISLABELED DATA

- * DL algorithms are quite robust to random errors in the training set
less robust to systemic errors

- * is it worth spending time in reducing the number of mislabeled images?

It depends! You should look at



10%	2%
0.6%	0.6%
9.4%	1.4%

FIRST YOU SHOULD TACKLE
YOUR OVERALL DEV SET ERROR

ONCE YOU SOLVE YOUR
DEV SET AND OTHER CAUSES
ERRORS, YOU SHOULD TACKLE
THE INCORRECT LABELS ISSUE.

- * correcting incorrect dev/test set examples:

- #1. apply same process to the dev and test sets (THEY NEED TO COME FROM THE SAME DISTRIBUTION)
- #2. examine examples your algorithm got right as well as those it got wrong.
- #3. train and dev/test set may now come from slightly \neq distributions. (NOT NECESSARILY BAD)

How to start prioritizing and see where to go next?

BUILD YOUR SYSTEM QUICKLY, THEN ITERATE!

MISMATCHED TRAINING AND DEV/TEST DATA

- * training and testing on \neq distributions:

#1 OPTION 1: combine both datasets, shuffle them, and split it into train/dev/test

advantage: same distribution / disadvantage: most data will come from one of the distributions.

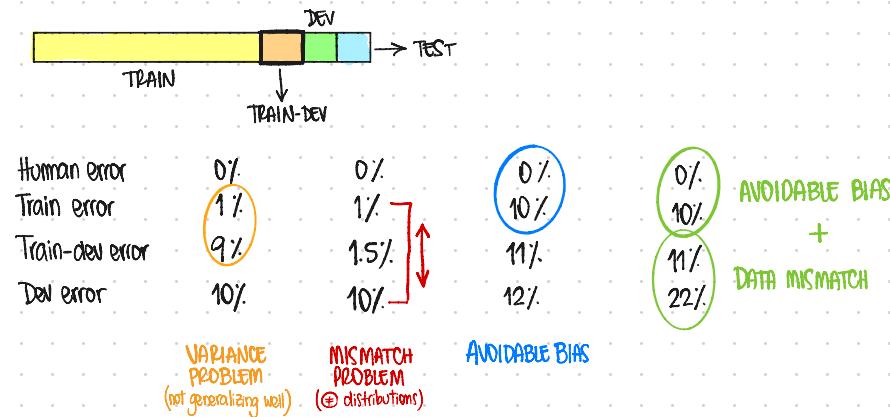
#2 OPTION 2: training set remains the same and dev/test sets will all come from a \neq distribution

advantage: you are aiming the target where you want it to be / disadvantage: train \neq dev/test

#3 OPTION 3: add 50% of new distribution to training set and dev/test have all the remaining

BIAS AND VARIANCE WITH MISMATCHED DATA DISTRIBUTIONS

- * training-dev set → same distribution as training set, but not used for training

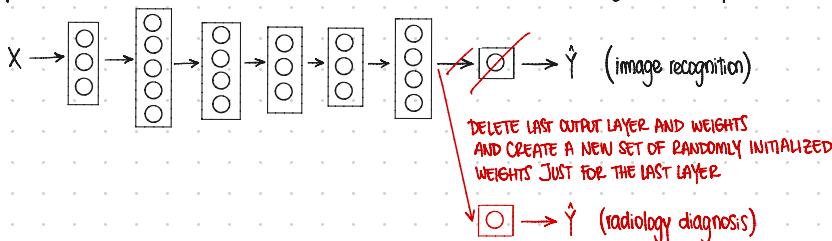


ADDRESSING DATA MISMATCH

- #1 Carry out manual error analysis to try to understand the difference between training dev/test
- #2 Make training data more similar or collect more data similar to dev/test set
 - * you can also do artificial data synthesis (but you might overfit)

TRANSFER LEARNING

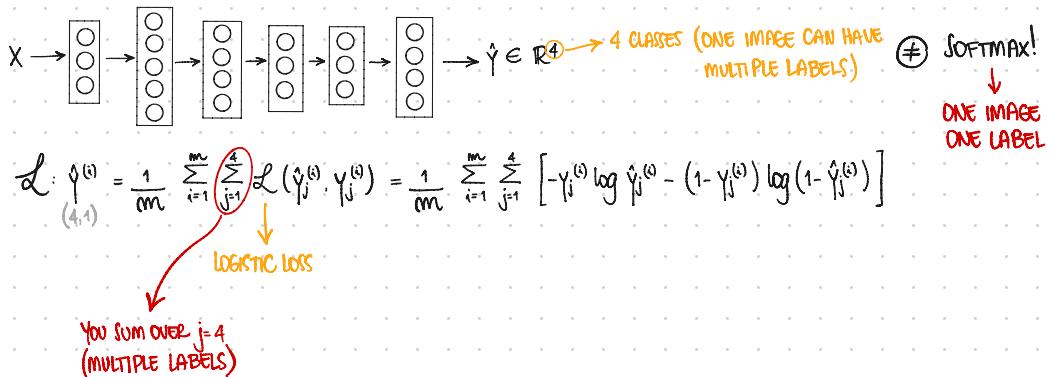
- * take the knowledge a NN learned from one task and apply that knowledge to a \oplus task.
- * You can train a NN to recognize cats and then use that knowledge to help you do a better job reading X-ray scans.



- * you swap the input of the NN and use a new one. So the NN has been pre-trained and you may fine-tune it.
- * when does transfer learning make sense? → when you have a lot of data from the problem you are transferring from and relatively less data for the problem you are transferring to.

MULTI-TASK LEARNING

- * You start off simultaneously trying to have one **NN** do several things at the same time. Each of these tasks helps hopefully the other tasks.



- * if you train a **NN** minimizing this cost function, then you are solving multiple problems at once. Since you are doing a multi-label classification, then you are saving time compared to having to train one **NN** for each type of label you have.

- * when does transfer learning make sense?

#1 Training on a set of tasks that could benefit from having shared lower-level features.

#2 amount of data you have for each task is quite similar (**NOT ALWAYS TRUE THO**)

#3 if you can train a big enough **NN** to do well on all tasks. The alternative is to train a **NN** for each task.

END-TO-END DEEP LEARNING

- * you may need a lot of data before you reach a decent accuracy.

- * used in face recognition and machine translation.



* instead of just take a whole image as input, the **NN** breaks it down into different steps to focus on the object you want to classify.

- * why does it work?

#1 the tasks are much simpler to solve when broken down into **#** tasks.

#2 there is usually a lot of data for all sub-tasks.

PROS

- * it lets the data speak
- * less hand-designing of components needed

CONS

- * may need large amounts of data
- * excludes potentially useful hand-designed components.